

# Week 2: Data Preparation

Visual Data Analytics

University of Sydney



# Outline

- Data frames
- Operations on data frames
- Data types
- Data profiling

# Data frames

# What is data?

- Data can be *structured* or *unstructured*
- Structured data has
  - Observations/ cases in rows
  - Variables/ features in columns
- Think of structured data as something that can be put into a spreadsheet.

# Example

	A	B	C	D	E
1	Player	Games Played	Minutes	Points	Birthplace
2	LeBron James	1394	53142	37860	Akron, USA
3	Kevin Durant	974	35776	26565	Washington DC, USA
4	James Harden	962	33378	23915	Los Angeles, USA
5	Russell Westbrook	1054	36352	23781	Long Beach, USA
6	Chris Paul	1178	40623	21235	Winston-Salem, USA
7	Stephen Curry	852	29256	20843	Akron, USA
8	DeMar DeRozan	993	34087	20812	Compton, USA
9	Damian Lillard	734	26648	18149	Oakland, USA
10	Rudy Gay	1089	34135	17463	New York, USA
11	Paul George	765	25711	15713	Palmdale, USA
12	Kyle Lowry	1055	33750	15556	Philadelphia, USA
13	Giannis Antetokounmpo	686	22365	15283	Athens, Greece
14	Anthony Davis	629	21641	15075	Chicago, USA
15	Bradley Beal	668	23211	14772	St Louis, USA
16	Kyrie Irving	637	21756	14770	Melbourne, Australia

Current as of January 2023

# US Cities and States

	A	B
1	Akron	Ohio
2	Los Angeles	California
3	Long Beach	California
4	Winston-Salem	North Carolina
5	Compton	California
6	Oakland	California
7	New York	New York
8	<u>Palmdale</u>	California
9	Philadelphia	Pennsylvania
10	Chicago	Illinois
11	St Louis	Missouri
12		

# Observations and variables

- In the first table:
  - What are the observations?
  - What are the variables?
- How about the second table?

# Pandas

- The Python library *pandas* stores data in a basic object called a *data frame*.
- It also provides functions for reading in and manipulating data frames.





# NBA Data

```
import pandas as pd
NBA = pd.read_csv('../data/NBA1.csv')
NBA
```

##		Player	Games Played	Minutes	Points	Birthpl
## 0		LeBron James	1394	53142	37860	Akron,
## 1		Kevin Durant	974	35776	26565	Washington DC,
## 2		James Harden	962	33378	23915	Los Angeles,
## 3		Russell Westbrook	1054	36352	23781	Long Beach,
## 4		Chris Paul	1178	40623	21235	Winston-Salem,
## 5		Stephen Curry	852	29256	20843	Akron,
## 6		DeMar DeRozan	993	34087	20812	Compton,
## 7		Damian Lillard	734	26648	18149	Oakland,
## 8		Rudy Gay	1089	34135	17463	New York,
## 9		Paul George	765	25711	15713	Palmdale,
## 10		Kyle Lowry	1055	33750	15556	Philadelphia,
## 11	Giannis	Antetokounmpo	686	22365	15283	Athens, Gre
## 12		Anthony Davis	629	21641	15075	Chicago,
## 13		Bradley Beal	668	23211	14772	St Louis,
## 14		Kyrie Irving	637	21756	14770	Melbourne, Austr

# Keys

- A key uniquely identifies an entry into the database.
- In the first table the player name can act as a key.
- Player birthplace could not be a key since both James and Curry are born in Akron.
- However, be careful with names as keys, since two players may have the same name.

# Atomicity

- Beware of variables that combine multiple pieces of information.
- Birthplace can be broken down into city and country.
- This may be useful if we want to visualise points scored by players born in USA and players born outside USA.
- Variables should be *atomic*
- This is also known as *parsing* the data.

# Split

```
new = NBA.Birthplace.str.split(',', expand = True)
new
```

##	0	1
## 0	Akron	USA
## 1	Washington DC	USA
## 2	Los Angeles	USA
## 3	Long Beach	USA
## 4	Winston-Salem	USA
## 5	Akron	USA
## 6	Compton	USA
## 7	Oakland	USA
## 8	New York	USA
## 9	Palmdale	USA
## 10	Philadelphia	USA
## 11	Athens	Greece
## 12	Chicago	USA
## 13	St Louis	USA
## 14	Melbourne	Australia

# With original table

```
NBA["CityOfBirth"] = new[0]
NBA["CountryOfBirth"] = new[1]
NBA
```

##	Player	Games Played	...	CityOfBirth	CountryOfBirth
## 0	LeBron James	1394	...	Akron	USA
## 1	Kevin Durant	974	...	Washington DC	USA
## 2	James Harden	962	...	Los Angeles	USA
## 3	Russell Westbrook	1054	...	Long Beach	USA
## 4	Chris Paul	1178	...	Winston-Salem	USA
## 5	Stephen Curry	852	...	Akron	USA
## 6	DeMar DeRozan	993	...	Compton	USA
## 7	Damian Lillard	734	...	Oakland	USA
## 8	Rudy Gay	1089	...	New York	USA
## 9	Paul George	765	...	Palmdale	USA
## 10	Kyle Lowry	1055	...	Philadelphia	USA
## 11	Giannis Antetokounmpo	686	...	Athens	Greece
## 12	Anthony Davis	629	...	Chicago	USA
## 13	Bradley Beal	668	...	St Louis	USA
## 14	Kyrie Irving	637	...	Melbourne	Australia

# Another example

---

Player	Points
Steph Curry	20843 for Warriors
Kevin Durant	17566 for Thunder, 5374 for Warriors, 3625 for Nets

---

# Split

Player	Points for Warriors	Points for Thunder	Points for Nets
Steph Curry	20843		
Kevin Durant	5374	17566	3625

# Problems

- Empty cells for Steph Curry
  - Not such a big issue (see missing data later)
- What if Kevin Durant moves to another team? (He did...)
- What if we want to include Giannis (only played for Bucks)?
  - Would need new columns and code that previously worked may break.



# Solution

Player	Team	Points
Steph Curry	Warriors	20843
Kevin Durant	Thunder	17566
Kevin Durant	Warriors	5374
Kevin Durant	Nets	3625

# First Normal Form

- New data entries can be added by adding rows only.
- Now the player and team combined form key.
- Overall the example so far is about getting the data into the *first normal form*.
- For the purposes of visualisation the lesson is to think carefully about how the data is structured.

# Operations on Data Frames

# The simple machines

- Hundreds of years ago, it was believed that all machines were made up of six simple machines
- These include: levers, wheels, pulleys, screws, etc.
- Nowadays machines are more complicated.
- But this is a good metaphor for data frames

# Simple machines of data

- We will consider six "simple machines" of data frames.
  - Transforming
  - Sorting
  - Filtering
  - Group by/ aggregate
  - Reshaping (melting and casting)
  - Joining (merging)
- By some combination of these we can 'munge' data frames into almost any data frame we need.

# Transform

- Create a new variable based on values of existing variables.
- For example, in the NBA data frame we have games played and points.
- Suppose we want to create a variable of points per game (PPG)

# Transform in Python

```
NBA["PPG"]=NBA["Points"]/ NBA["Games Played"]
NBA
```

##	Player	Games Played	...	CountryOfBirth	PPG
## 0	LeBron James	1394	...	USA	27.159254
## 1	Kevin Durant	974	...	USA	27.274127
## 2	James Harden	962	...	USA	24.859667
## 3	Russell Westbrook	1054	...	USA	22.562619
## 4	Chris Paul	1178	...	USA	18.026316
## 5	Stephen Curry	852	...	USA	24.463615
## 6	DeMar DeRozan	993	...	USA	20.958711
## 7	Damian Lillard	734	...	USA	24.726158
## 8	Rudy Gay	1089	...	USA	16.035813
## 9	Paul George	765	...	USA	20.539869
## 10	Kyle Lowry	1055	...	USA	14.745024
## 11	Giannis Antetokounmpo	686	...	Greece	22.278426
## 12	Anthony Davis	629	...	USA	23.966614
## 13	Bradley Beal	668	...	USA	22.113772
## 14	Kyrie Irving	637	...	Australia	23.186813
##					

# Sort

- Suppose we want to sort the data according to one of the variables.
- We can use the `sort_values` function.
- Consider that we want to sort by minutes played from smallest to largest.



# Players by minutes

```
NBAbymin = NBA.sort_values(by = 'Minutes')
NBAbymin
```

##	Player	Games Played	...	CountryOfBirth	PPG
## 12	Anthony Davis	629	...	USA	23.966614
## 14	Kyrie Irving	637	...	Australia	23.186813
## 11	Giannis Antetokounmpo	686	...	Greece	22.278426
## 13	Bradley Beal	668	...	USA	22.113772
## 9	Paul George	765	...	USA	20.539869
## 7	Damian Lillard	734	...	USA	24.726158
## 5	Stephen Curry	852	...	USA	24.463615
## 2	James Harden	962	...	USA	24.859667
## 10	Kyle Lowry	1055	...	USA	14.745024
## 6	DeMar DeRozan	993	...	USA	20.958711
## 8	Rudy Gay	1089	...	USA	16.035813
## 1	Kevin Durant	974	...	USA	27.274127
## 3	Russell Westbrook	1054	...	USA	22.562619
## 4	Chris Paul	1178	...	USA	18.026316
## 0	LeBron James	1394	...	USA	27.159254
##					
## 115	...	...	...	...	...

# Filter

- Filtering involves selecting only some subset of the data.
- There are many ways to do this
  - Select rows
  - Select columns
- Select by a logical condition

# Example

- Suppose we only want to consider
  - Players with points per game greater than 20
  - Players born in Akron
  - Players not born in the United States
- These are all examples of logical conditions (either true or false).

# Players with PPG above 20

```
NBAppg20 = NBA.loc[NBA["PPG"]>20]
NBAppg20
```

```
##           Player  Games Played  ...  CountryOfBirth      PPG
## 0      LeBron James      1394  ...           USA  27.159254
## 1      Kevin Durant      974  ...           USA  27.274127
## 2      James Harden      962  ...           USA  24.859667
## 3  Russell Westbrook     1054  ...           USA  22.562619
## 5      Stephen Curry      852  ...           USA  24.463615
## 6      DeMar DeRozan      993  ...           USA  20.958711
## 7      Damian Lillard      734  ...           USA  24.726158
## 9      Paul George       765  ...           USA  20.539869
## 11  Giannis Antetokounmpo   686  ...         Greece  22.278426
## 12      Anthony Davis      629  ...           USA  23.966614
## 13      Bradley Beal       668  ...           USA  22.113772
## 14      Kyrie Irving       637  ...        Australia  23.186813
##
## [12 rows x 8 columns]
```

# Players born in Akron

```
NBAAkr = NBA.loc[NBA["CityOfBirth"] == 'Akron']  
NBAAkr
```

```
##           Player  Games Played  Minutes  ...  CityOfBirth CountryOfBirth  
## 0    LeBron James           1394    53142  ...      Akron             USA 2  
## 5    Stephen Curry            852    29256  ...      Akron             USA 2  
##  
## [2 rows x 8 columns]
```

Note that a single = denotes assignment, a double == denotes 'equals' in a logical statement.

# Players born outside USA

```
NBAnonUS = NBA.loc[NBA["CountryOfBirth"] != 'USA']  
NBAnonUS
```

```
##           Player  Games Played  ...  CountryOfBirth      PPG  
## 11  Giannis Antetokounmpo      686  ...           Greece  22.278426  
## 14      Kyrie Irving      637  ...        Australia  23.186813  
##  
## [2 rows x 8 columns]
```

In general we can read ! as 'not' in Python

# Group by / aggregate

- Suppose we want to compare total points scored by players country of birth.
- This requires two functions
  - The `groupby` function tells us the variable to group on (in this case Country of Birth).
  - The `agg` function tells us which variable to aggregate (in this case points)

# Groupby/aggregate in Python

```
NBAg = NBA.groupby('CountryOfBirth').agg({'Points': 'sum'})  
NBAg
```

```
##                Points  
## CountryOfBirth  
## Australia      14770  
## Greece         15283  
## USA            271739
```

Other ways to aggregate include mean, min and max.



# Reshape

- Often in order to produce the visualisation we want we need to reshape the data.
- This is done using two functions
  - The function `melt` converts the data from wide to long.
  - The function `pivot` converts the data from long to wide.

# Melting

```
NBAlong = NBA.melt(id_vars=['Player'],value_vars=['Games Played', 'Minutes Played'])
NBAlong
```

##	Player	variable	value
## 0	LeBron James	Games Played	1394
## 1	Kevin Durant	Games Played	974
## 2	James Harden	Games Played	962
## 3	Russell Westbrook	Games Played	1054
## 4	Chris Paul	Games Played	1178
## 5	Stephen Curry	Games Played	852
## 6	DeMar DeRozan	Games Played	993
## 7	Damian Lillard	Games Played	734
## 8	Rudy Gay	Games Played	1089
## 9	Paul George	Games Played	765
## 10	Kyle Lowry	Games Played	1055
## 11	Giannis Antetokounmpo	Games Played	686
## 12	Anthony Davis	Games Played	629
## 13	Bradley Beal	Games Played	668
## 14	Kyrie Irving	Games Played	637
## 15	LeBron James	Minutes	53142

# Pivoting

```
NBAwide = NBAlong.pivot(index='Player', columns = 'variable')
NBAwide
```

##	value
## variable	Games Played Minutes Points
## Player	
## Anthony Davis	629 21641 15075
## Bradley Beal	668 23211 14772
## Chris Paul	1178 40623 21235
## Damian Lillard	734 26648 18149
## DeMar DeRozan	993 34087 20812
## Giannis Antetokounmpo	686 22365 15283
## James Harden	962 33378 23915
## Kevin Durant	974 35776 26565
## Kyle Lowry	1055 33750 15556
## Kyrie Irving	637 21756 14770
## LeBron James	1394 53142 37860
## Paul George	765 25711 15713
## Rudy Gay	1089 34135 17463
## Russell Westbrook	1054 36352 23781
## Stephen Curry	853 28256 20242

# A better example

```
Sydney = pd.read_csv('../data/SydneyClimate.csv')  
Sydney
```

##	Year	M01	M02	M03	M04	M05	M06	M07	M08	M09	M10	M11
## 0	2017	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	24.9	24.8
## 1	2018	28.5	28.1	27.4	26.7	22.2	17.7	19.9	19.3	21.0	21.9	25.0
## 2	2019	29.6	27.7	26.9	25.1	22.7	18.6	19.8	19.5	22.0	24.7	27.0
## 3	2020	29.0	27.5	25.6	24.5	20.3	18.7	18.2	19.5	22.7	24.2	26.1
## 4	2021	27.4	26.7	25.5	24.3	20.9	18.1	18.2	20.7	22.5	24.2	23.2
## 5	2022	27.7	26.7	24.8	23.7	20.8	18.3	17.0	19.8	20.6	22.6	24.3


Retrieved from Bureau of Meteorology

# Melting

```
Sydlong = Sydney.melt(id_vars='Year').sort_values(by = ['Year', 'variable'])  
Sydlong
```

```
##      Year variable  value  
##  0    2017      M01    NaN  
##  6    2017      M02    NaN  
## 12    2017      M03    NaN  
## 18    2017      M04    NaN  
## 24    2017      M05    NaN  
## ..     ...      ...     ...  
## 47    2022      M08    19.8  
## 53    2022      M09    20.6  
## 59    2022      M10    22.6  
## 65    2022      M11    24.3  
## 71    2022      M12     NaN  
##  
## [72 rows x 3 columns]
```

# Plot



# Merge

- Bring two data frames together
- Similar to a VLOOKUP type function in spreadsheet programs such as Excel.
- We can use `merge` to add information about State of birth for NBA players.

# Merge

```
CitiesStates = pd.read_csv('../data/UScitiesstates.csv')
NBAMerge = pd.merge(NBA,CitiesStates, left_on = 'CityOfBirth', right_on =
NBAMerge
```

```
##           Player  Games Played  ...           City           State
## 0      LeBron James          1394  ...           Akron           Ohio
## 1      Stephen Curry           852  ...           Akron           Ohio
## 2      James Harden           962  ...  Los Angeles  California
## 3  Russell Westbrook          1054  ...   Long Beach  California
## 4      Chris Paul          1178  ... Winston-Salem  North Carolina
## 5      DeMar DeRozan           993  ...      Compton  California
## 6      Damian Lillard           734  ...      Oakland  California
## 7      Rudy Gay          1089  ...    New York    New York
## 8      Paul George           765  ...    Palmdale  California
## 9      Kyle Lowry          1055  ... Philadelphia  Pennsylvania
## 10     Anthony Davis           629  ...    Chicago    Illinois
## 11     Bradley Beal           668  ...    St Louis    Missouri
##
## [12 rows x 10 columns]
```



# Different type of merge

- Left: Keep all entries from first data frame
- Right: Keep all entries from second data frame
- Inner: Keep all entries that appear in both data frames
- Outer: Keep all entries that appear in either data frame

# Merge

```
CitiesStates = pd.read_csv('../data/UScitiesstates.csv')
NBAMerge = pd.merge(NBA,CitiesStates, how = 'outer', left_on = 'CityOfB
NBAMerge
```

##	Player	Games Played	...	City	State
## 0	LeBron James	1394	...	Akron	Ohio
## 1	Stephen Curry	852	...	Akron	Ohio
## 2	Kevin Durant	974	...	NaN	NaN
## 3	James Harden	962	...	Los Angeles	California
## 4	Russell Westbrook	1054	...	Long Beach	California
## 5	Chris Paul	1178	...	Winston-Salem	North Carolina
## 6	DeMar DeRozan	993	...	Compton	California
## 7	Damian Lillard	734	...	Oakland	California
## 8	Rudy Gay	1089	...	New York	New York
## 9	Paul George	765	...	Palmdale	California
## 10	Kyle Lowry	1055	...	Philadelphia	Pennsylvania
## 11	Giannis Antetokounmpo	686	...	NaN	NaN
## 12	Anthony Davis	629	...	Chicago	Illinois
## 13	Bradley Beal	668	...	St Louis	Missouri
## 14	Kyrie Irving	637	...	NaN	NaN

# Putting them together

- In your own time, construct data frames for the following:
  - Each observation is a state and with the maximum points per minute (PPM) by a player from each state.
  - The same as above with states ranked from highest to lowest according to the maximum PPM.
  - The same as above but with an extra column with the player name of the player with the highest PPM in each state.
- There may be more than one correct answer.

# Data types

# Data Types

- Each variable measures a certain characteristic.
- Characteristics can be measured in different ways
- This leads to *data type* which are important for understanding
  - How we can transform data.
  - The correct visualisation to use.

# Scales of measurement

- Any old (or new) statistics textbook will introduce four scales of measurement
  - Nominal
  - Ordinal
  - Interval
  - Ratio
- These are still useful (with some caveats).

# Nominal data

- Tells us something about a characteristic but there is no notion of having more or less of a characteristic.
- Example: Country of birth.
- Can you think of other examples?
- Even if we assign numbers to nominal categories, it does not make sense to find means medians etc.
- The mode still makes sense.

# Ordinal data

- Tells us whether we have more or less of a characteristic, but not how much more or less.
- Example: rate players as good, very good, excellent.
- If we assign numbers to nominal categories it still does not make sense to add or subtract these numbers.
- However the median (and the mode) still make sense.



# Interval/Ratio data

- All numerical data is either interval or ratio data.
- The differences between the two concern whether the zero point of the scale truly represents an absence of the characteristic being measured.
- Best understood with an example.

# Points per game

- Suppose I constructed a new index for points per game (PPG) where a PPG of 20 now becomes a PPG of 0.
  - Paul George (PPG: 20.5) would have a "new" PPG of 0.5.
  - Steph Curry (PPG: 24.5) would have a "new" PPG of 4.5.
- Does this mean that Curry is scoring 9 times as much as George? No.

# The textbook example

- The famous example is temperature.
- The Celsius scale attaches 0 and 100 to the freezing and boiling point of water.
  - This is arbitrary
- For the Kelvin scale, zero is true zero since it is a temperature where atoms have no energy (loosely speaking).

# Does it matter?

- Not that much (outside of science).
- Most data we see in business are ratio data.
- In general, I will use *numeric data* and *ratio data* interchangeably.
- Just think carefully when dividing with numerical variables.

# Summary

Operation	Nominal	Ordinal	Interval	Ratio
Equality	✓	✓	✓	✓
Order		✓	✓	✓
Add / subtract			✓	✓
Multiply / divide				✓
Mode	✓	✓	✓	✓
Median		✓	✓	✓
Arithmetic mean			✓	✓
Geometric mean				✓

# Some exceptions

- Nominal data with two categories.
  - Born in US assigned 1, born outside US assigned 0.
  - Arithmetic mean is then the *proportion* born in US.
- Likert (customer satisfaction) scales:
  - Strongly disagree = 1, Disagree = 2, etc.
  - Using an arithmetic mean is controversial but common in practice.
- Time is very unusual since calendar effects are important in business.

# Types in Pandas

```
NBA.dtypes
```

```
## Player          object
## Games Played    int64
## Minutes         int64
## Points          int64
## Birthplace      object
## CityOfBirth     object
## CountryOfBirth  object
## PPG             float64
## dtype: object
```

Object is text, int64 is an integer and float64 is a real number.

# Data Profiling



# Some issues

- Duplicated entries
  - Can be removed using `drop_duplicates` function in pandas.
- Data entry errors
  - For example, one cannot score 20000 points in 3 minutes.
  - Requires domain knowledge.
  - These issues can be discovered during visualisation.

# Standardisations

- Steph Curry may appear elsewhere in the data as
  - "Stephen Curry"
  - "Wardell Stephen Curry II" (his full name).
- Similar things happen with company names  
Facebook/ Meta, General Motors/ GM etc.
- Requires domain knowledge and some extra coding.

# Missing Data

- There are often missing data
- Need to think of a good strategy to encode missing data.
- In Python there is NaN for this.
- Often will replace with a number but this is a bad strategy


# Do not use "0" for "missing"

- Suppose if games played is missing and we replace with zero.
- The zeros will distort the mean.
- If we compute points per game there will be division by zero.
- Do not replace missing values by zeros.

# Do not use "-999" for "missing"

- Sometimes a completely implausible number such as -999 is used to denote missing.
- This can lead to strange visualisations
- The following example is for the Sydney temperature data

# Temperature



# Dealing with missing data

- Only use complete cases
- Impute missing values
  - With a random value
  - With mean/median or mode
  - More complicated models
- Report/Visualise missing data
  - By reporting missing data this gives a better idea of uncertainty.

# Report missing

**Do you approve or disapprove of the job Anthony Albanese is doing as Prime Minister?**

	◀ TOTAL ▶	◀ Labor ▶	◀ TOTAL: Coalition ▶	◀ Greens ▶	◀ Minor parties/Independents ▶
Strongly approve	18%	30%	10%	22%	9%
Approve	42%	58%	31%	49%	29%
Disapprove	16%	5%	26%	13%	23%
Strongly disapprove	12%	0%	22%	2%	26%
Don't know	13%	7%	11%	14%	13%
TOTAL: Approve	60%	88%	41%	71%	39%
TOTAL: Disapprove	28%	5%	48%	15%	49%

Source: **Guardian**



# Wrap-up

# Conclusions

- Our focus from now on will be visualisation.
- Writing code to visualise messy data is hard.
- Spend the time to clean your data.
- Focus on the *principles* since these work for Excel, Tableau, R, Python, etc.
- After you appreciate the principles, practice your coding.

# Questions