

# Week 3: Visualising One Variable

Visual Data Analytics

University of Sydney



# Outline

- Nominal/Ordinal Data
  - Bar
  - Lollipop
  - Pie/donut
- Numeric data
  - Box plot
  - Histograms
  - Kernel density

# Motivation

- Understand the distribution of a variable
  - Find outliers
  - Find multi-modality
  - Find skew
- Understanding the distribution is about generating interesting questions for further analysis.
- Thinking probabilistically is about thinking about distributions and not just the mean.

# Examples

- We will use two datasets that can be directly loaded from the `seaborn` package.
  - The `taxis` dataset with data on pickup and drop off locations, fares, payment type etc., in New York City.
  - The `diamonds` dataset with information on size, cut clarity, price, etc. of diamonds.
- These contain categorical (nominal and ordinal) and numeric variables.

# Categorical variables

# The bar chart

- Categories displayed on one axis (usually x).
- The *frequency* of each observation is displayed on the other axis (usually y).
- The frequency is mapped to the *length* of each bar.
- For this reason always include zero on the y axis.


# Taxis data

```
import seaborn as sns
taxisdat = sns.load_dataset('taxi')
taxisdat
```

```
##           pickup           dropoff ... pickup_borough dropoff
## 0      2019-03-23 20:21:09  2019-03-23 20:27:24 ...      Manhattan
## 1      2019-03-04 16:11:55  2019-03-04 16:19:00 ...      Manhattan
## 2      2019-03-27 17:53:01  2019-03-27 18:00:25 ...      Manhattan
## 3      2019-03-10 01:23:59  2019-03-10 01:49:51 ...      Manhattan
## 4      2019-03-30 13:27:42  2019-03-30 13:37:14 ...      Manhattan
## ...           ...           ... ...           ...
## 6428  2019-03-31 09:51:53  2019-03-31 09:55:27 ...      Manhattan
## 6429  2019-03-31 17:38:00  2019-03-31 18:34:23 ...      Queens
## 6430  2019-03-23 22:55:18  2019-03-23 23:14:25 ...      Brooklyn
## 6431  2019-03-04 10:09:25  2019-03-04 10:14:29 ...      Brooklyn
## 6432  2019-03-13 19:31:22  2019-03-13 19:48:02 ...      Brooklyn
##
## [6433 rows x 14 columns]
```

# Bar plot of pick up borough


```
sns.countplot(data = taxisdat, x='pickup_borough')
```






# Change orientation

```
sns.countplot(data = taxisdat, y='pickup_borough')
```



# Order by frequency

```
sns.countplot(data = taxisdat, x='pickup_borough', order = taxisdat['pi
```



Data are nominal - this is fine.

# Ordinal data

- For nominal data it is suitable, to order according to frequency.
- This is not the case for ordinal data
- Always order according to categories of the variable.
- Diamonds dataset has clarity as an ordinal variable
  - Categories ordered as IF, VVS1, VVS2, VS1, VS2, SI1, SI2, I1.


# Diamonds data

```
diam = sns.load_dataset('diamonds')  
diam
```

```
##          carat          cut color clarity  depth  table  price         x         y  
## 0         0.23        Ideal     E    SI2   61.5   55.0    326    3.95    3.98    2.4  
## 1         0.21     Premium     E    SI1   59.8   61.0    326    3.89    3.84    2.3  
## 2         0.23         Good     E    VS1   56.9   65.0    327    4.05    4.07    2.3  
## 3         0.29     Premium     I    VS2   62.4   58.0    334    4.20    4.23    2.6  
## 4         0.31         Good     J    SI2   63.3   58.0    335    4.34    4.35    2.7  
## ...         ...         ...    ...    ...    ...    ...    ...    ...    ...  
## 53935      0.72        Ideal     D    SI1   60.8   57.0   2757    5.75    5.76    3.5  
## 53936      0.72         Good     D    SI1   63.1   55.0   2757    5.69    5.75    3.6  
## 53937      0.70    Very Good     D    SI1   62.8   60.0   2757    5.66    5.68    3.5  
## 53938      0.86     Premium     H    SI2   61.0   58.0   2757    6.15    6.12    3.7  
## 53939      0.75        Ideal     D    SI2   62.2   55.0   2757    5.83    5.87    3.6  
##  
## [53940 rows x 10 columns]
```


# Ordinal

```
diam = sns.load_dataset('diamonds')  
sns.countplot(data=diam, x='clarity')
```




Categories ordered by levels of variable - this is fine.

# Incorrect plot



Incorrect. Categories in alphabetical order.


# Incorrect plot



Incorrect. Ordered by frequency.

# Single color

```
diam = sns.load_dataset('diamonds')  
sns.countplot(data=diam, x='clarity', color='tab:blue')
```





# Coloring

- Although by default categories have different colors this is not strictly necessary.
- Arguably it is distracting, especially when there are more categories.
- Later on we will use color to display data
  - For example grouping by a second variable and mapping that to color.
- This will be covered later on.

# Lollipop charts

- If there are
  - A large number of categories,
  - If the categories all have similar frequencies,
- then consider using a lollipop chart.
- This can be done with some data munging using `value_counts` and the `stem` function in `matplotlib`.

# Data preparation

For simpler graph, will only consider dropoff in Manhattan


```
freq = taxisdat[taxisdat['pickup_borough']=='Bronx'].value_counts('pick  
freq
```

```
## pickup_zone  
## East Concourse/Concourse Village      9  
## Parkchester                           8  
## Mott Haven/Port Morris                8  
## Co-Op City                            5  
## East Tremont                          5  
## Mount Hope                            5  
## Claremont/Bathgate                    5  
## Soundview/Castle Hill                 4  
## Morrisania/Melrose                    4  
## Van Nest/Morris Park                  3  
## University Heights/Morris Heights     3
```

# Lollipop plot (code)

```
import matplotlib.pyplot as plt
plt.stem(freq)
plt.xticks(range(1, len(freq.index)+1), freq.index, rotation='vertical')
plt.show()
```




# Lollipop plot (output)






# Pie charts

- Pie charts are considered to be poor practice by visualisation experts since
  - It is difficult to compare sizes of angles.
  - It is difficult to make comparisons unless categories are close.
  - They do not handle large numbers of categories.
- Following examples come from a **Business Insider** article by Walt Hickey.

# Pie chart




# Bar chart






# Pie chart



# Bar chart



# How to do pie charts

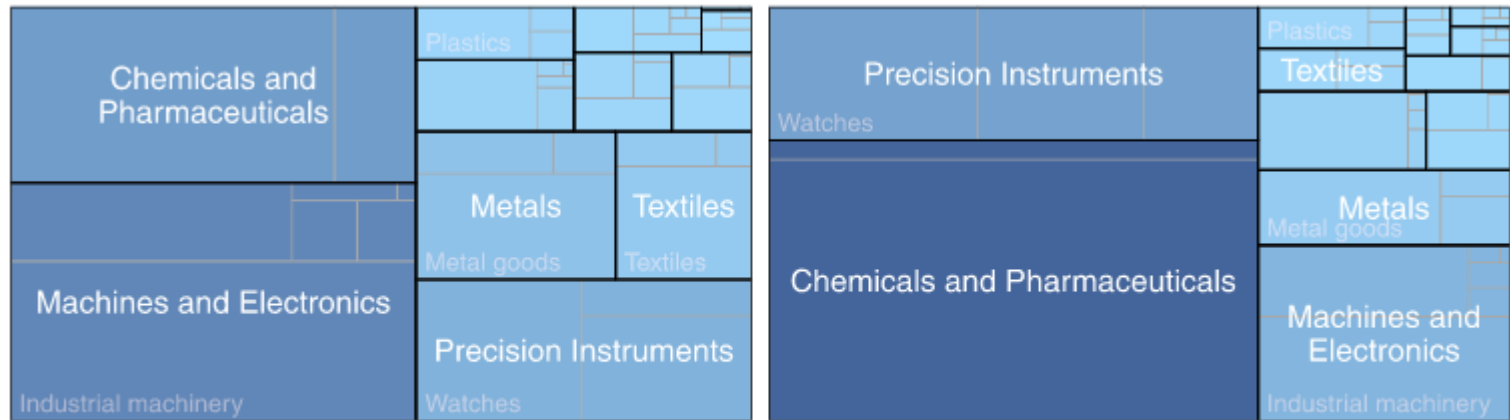
- If you absolutely MUST do a pie chart a guide can be found at this [link](#).
- A donut chart is a pie chart with a hole. It is even worse than a pie chart.



# Treemaps

- Even bar charts can struggle when the number of categories is truly huge.
- One way to handle this is using a *treemap*.
- See [this example](#)
- These are particularly well suited when categories follow a hierarchy.
- The following example considers Swiss exports that are classified into 12 categories and 48 subcategories.

# Swiss Exports



Regional or Categorical Share of Exports (in %) 0 20 40 60


# Numerical Data

# Histogram

- The equivalent of a bar chart for numerical data is a histogram.
- The area of each bar represents the frequency within a certain interval.
- If all bars have equal width then frequency is mapped to the length of the bars too.
- Zero should always be included on the y axis (but not necessarily x axis).

# Histogram

```
sns.histplot(taxisdat[ 'fare' ])
```






# What do we see?

- Right skew
  - Should we use mean or median as measure of central tendency?
- A few big outliers.
- A spike (second 'mode') at around \$50
  - Could represent a fixed fee (e.g. from airport).


# Change number of bins

```
sns.histplot(taxisdat['fare'], bins=10)
```



# Change number of bins

```
sns.histplot(taxisdat['fare'], bins=2000)
```



# Lessons

- By having too many (or too few) bins we can miss out on important features of data.
- In the above example the spike of fares around \$50 is not seen when the number of bins is changed.
- In general default choice of bin number is good, however it is always a good idea to experiment.

# Kernel density estimate (KDE)

- A kernel density estimates the *probability density function (pdf)* of data.
- For data  $x_1, x_2, \dots, x_n$  the KDE is given by


$$\hat{f}(x) = \frac{1}{n} \sum_{i=1}^n K_h(x - x_i)$$

- The function  $K_h(\cdot)$  is called the *kernel*.
- Can take many forms.
- The function depends on a bandwidth  $h$  (to be explained soon).


# Intuition behind KDE

- If I observe a point in some location, that evidence supports that there is probability that a point comes from a nearby region.
- Imagine I drop a mountain of sand at the location I observe the data point.
- The shape of the sand is the kernel function.
- If I repeat this for  $n$  observations the result is the KDE.

# KDE (n-1)




# KDE (n=2)







# KDE (n=3)



# KDE (n=4)



# KDE (n=20)




# The bandwidth

- The bandwidth  $h$  controls whether the mountain of sand is 'peaked' or 'flat' .
- For small bandwidth the mountain of sand is more peaked and the KDE is more wiggly.
- For large bandwidth the mountain of sand is more flat and the KDE is more smooth.
- This is similar to the role of the number of bins in the histogram.
- There are sensible defaults used by visualisation packages.


# KDE plot

```
sns.kdeplot(taxisdat['fare'])
```




# KDE plot (double default BW)

```
sns.kdeplot(taxisdat['fare'], bw_adjust = 2)
```



# KDE plot (half default BW)

```
sns.kdeplot(taxisdat['fare'], bw_adjust = 0.5)
```




# Violin plot

- A violin plot mirrors a KDE and fills it in.
- It is particularly useful for making comparisons of density according to a grouping variable.
- We will cover this next week.



# Violin plot

```
sns.violinplot(taxisdat['fare'])
```




# Boxplot

- Inside the violin plot is a boxplot.
- The boxplot is a summary of five statistics
  - Median
  - First Quartile
  - Third Quartile
  - Minimum
  - Maximum

# Box plot

```
sns.boxplot(taxisdat['fare'])
```



# Fences

- For most implementations, a boxplot actually shows an upper and lower fence rather than the maximum and minimum.
- The upper (lower) fence is given by the third (first) quartile plus (minus) 1.5 times the IQR.
- The maximum (minimum) is shown instead if it is less (greater) than the upper (lower) fence.

# Boxplots v KDE


- Note that in this example the spike at around \$50 is lost in the boxplot.
- However it is clearer that there are four outliers above \$110.
- There is no right and wrong answer, it all depends on what you are trying to visualise.

# Rug plot

- The final plot we will consider is a rug plot.
- The rug plot can highlight outliers.
- It is harder to understand the shape of the distribution using a rug plot, especially for large sample sizes.
- As a univariate plot, a jittered rug plot (strip plot) works better.


# Rug plot

```
sns.rugplot(taxisdat['fare'])
```



# Rug plot (jittered)

```
sns.stripplot(y=taxisdat['fare'])
```





# Wrap-up

# Conclusions

- Univariate plots are useful for
  - Understanding distribution of a variable
  - Finding outliers
  - Finding frequent values
  - Seeing whether data are skewed.
- Always remember that univariate plots generate questions. To answer these questions requires domain knowledge and further analysis.

# Questions